



ECM2414 – Software Development

Concurrency I

Dr. Yulei Wu, T1:W2 - L1

Department of Computer Science
College of Engineering, Mathematics and Physical Sciences
University of Exeter



Today's lecture

- Concurrency
 - What is it? Why use it?
- Multi-threading
 - As separate from multi-tasking
- Creating, starting, and stopping a thread
- Issues with multi-threading

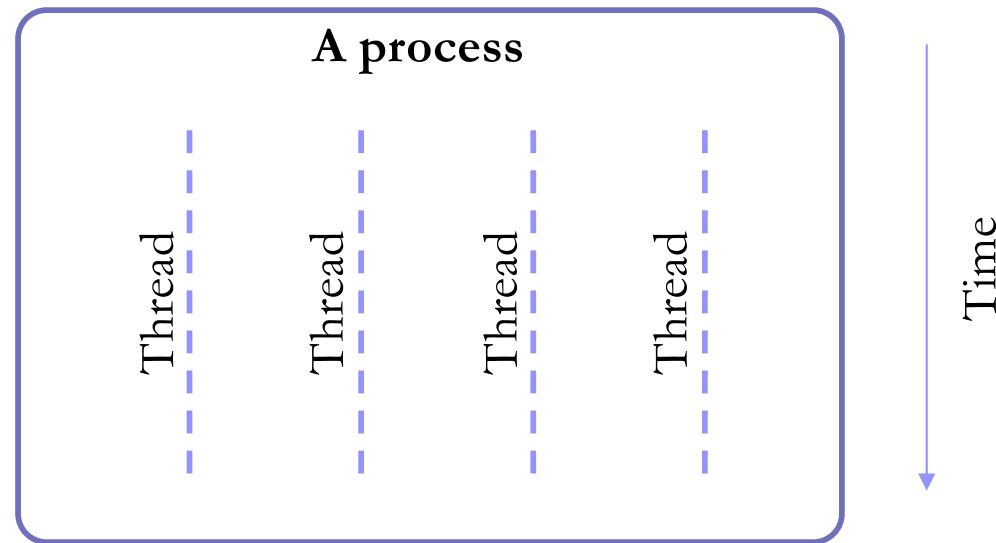


Concurrency

- To be able to perform multiple **processes** at the same time – including the potential interaction between them
 - Multiple tasks at the same time sometimes called ‘parallelism’
- Why are we interested in concurrency?
 - Allowing programs to interact with other systems/users
 - Multi-core systems

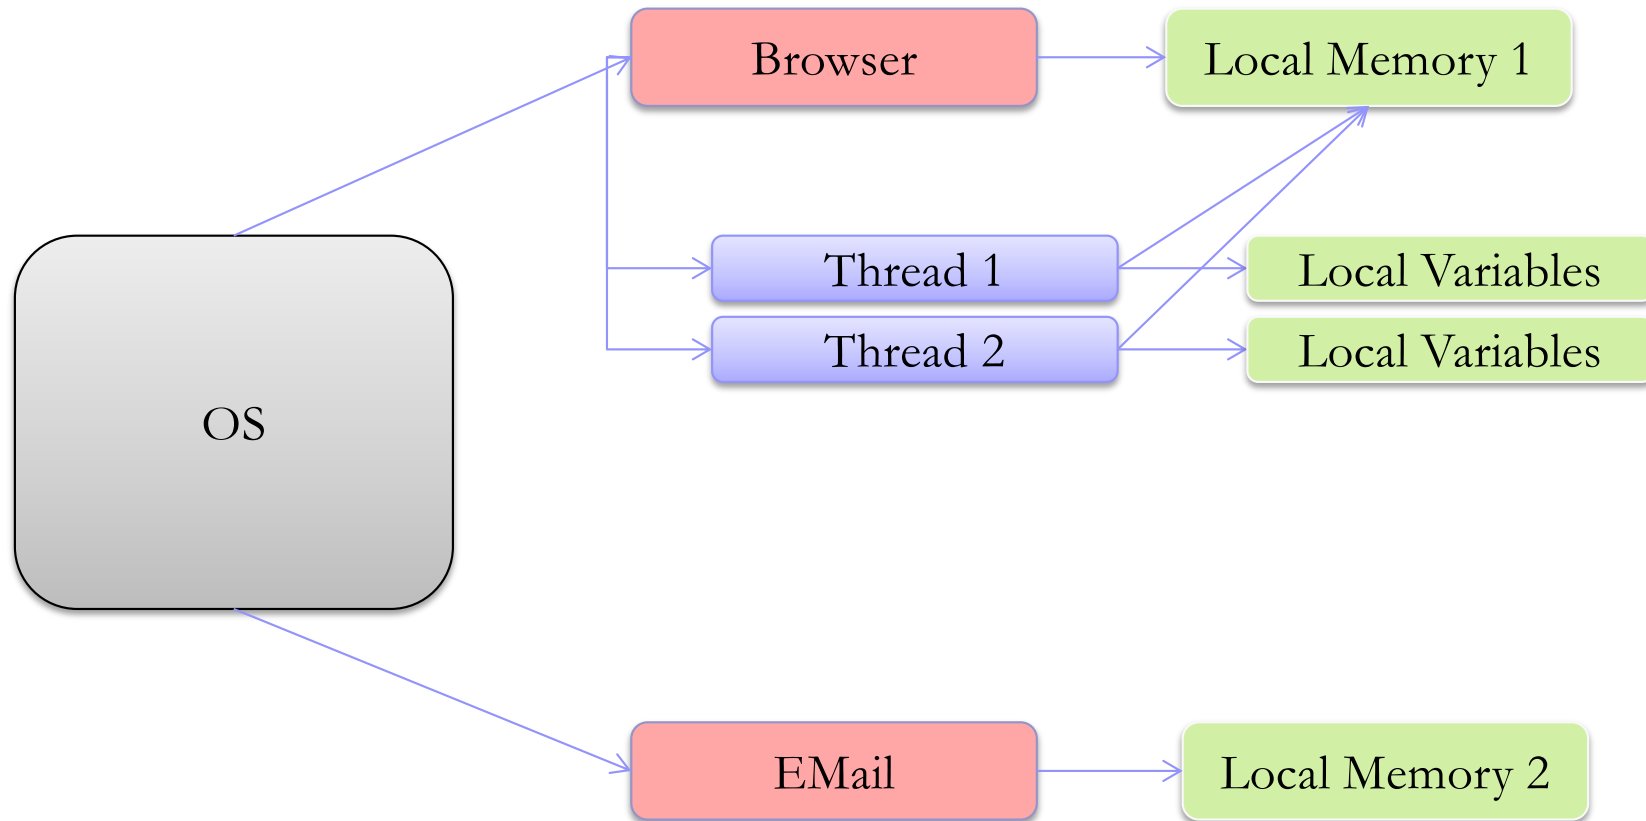
Threads

- Can be treated as lightweight processes.
- Threads exist within a process



- Why thread? Why not have separate programs?
 - Shared memory space

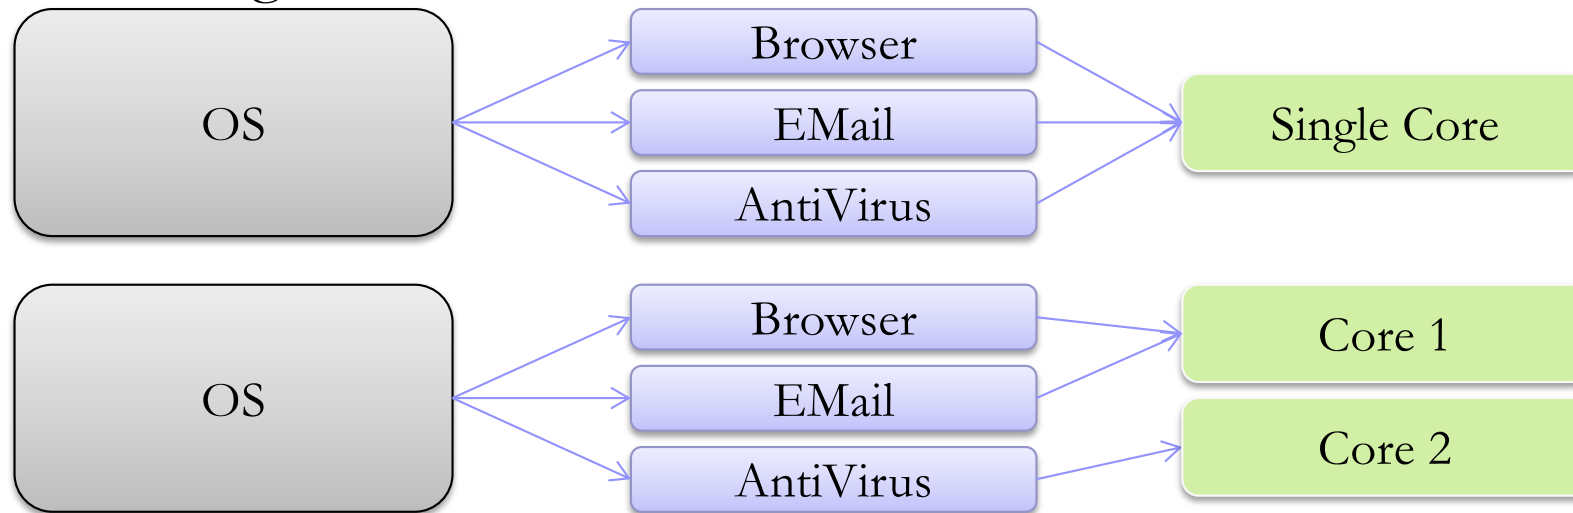
Multi-threading



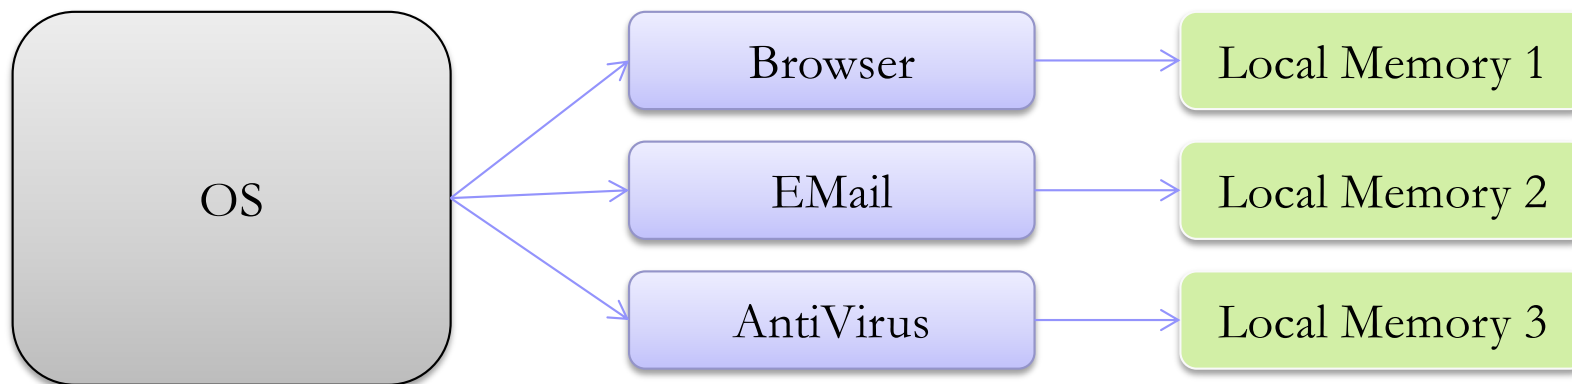
- The OS can run threads in parallel
 - Virtually, on a single processor
 - Actually on a multi-core processor

Multi-tasking

■ Processing

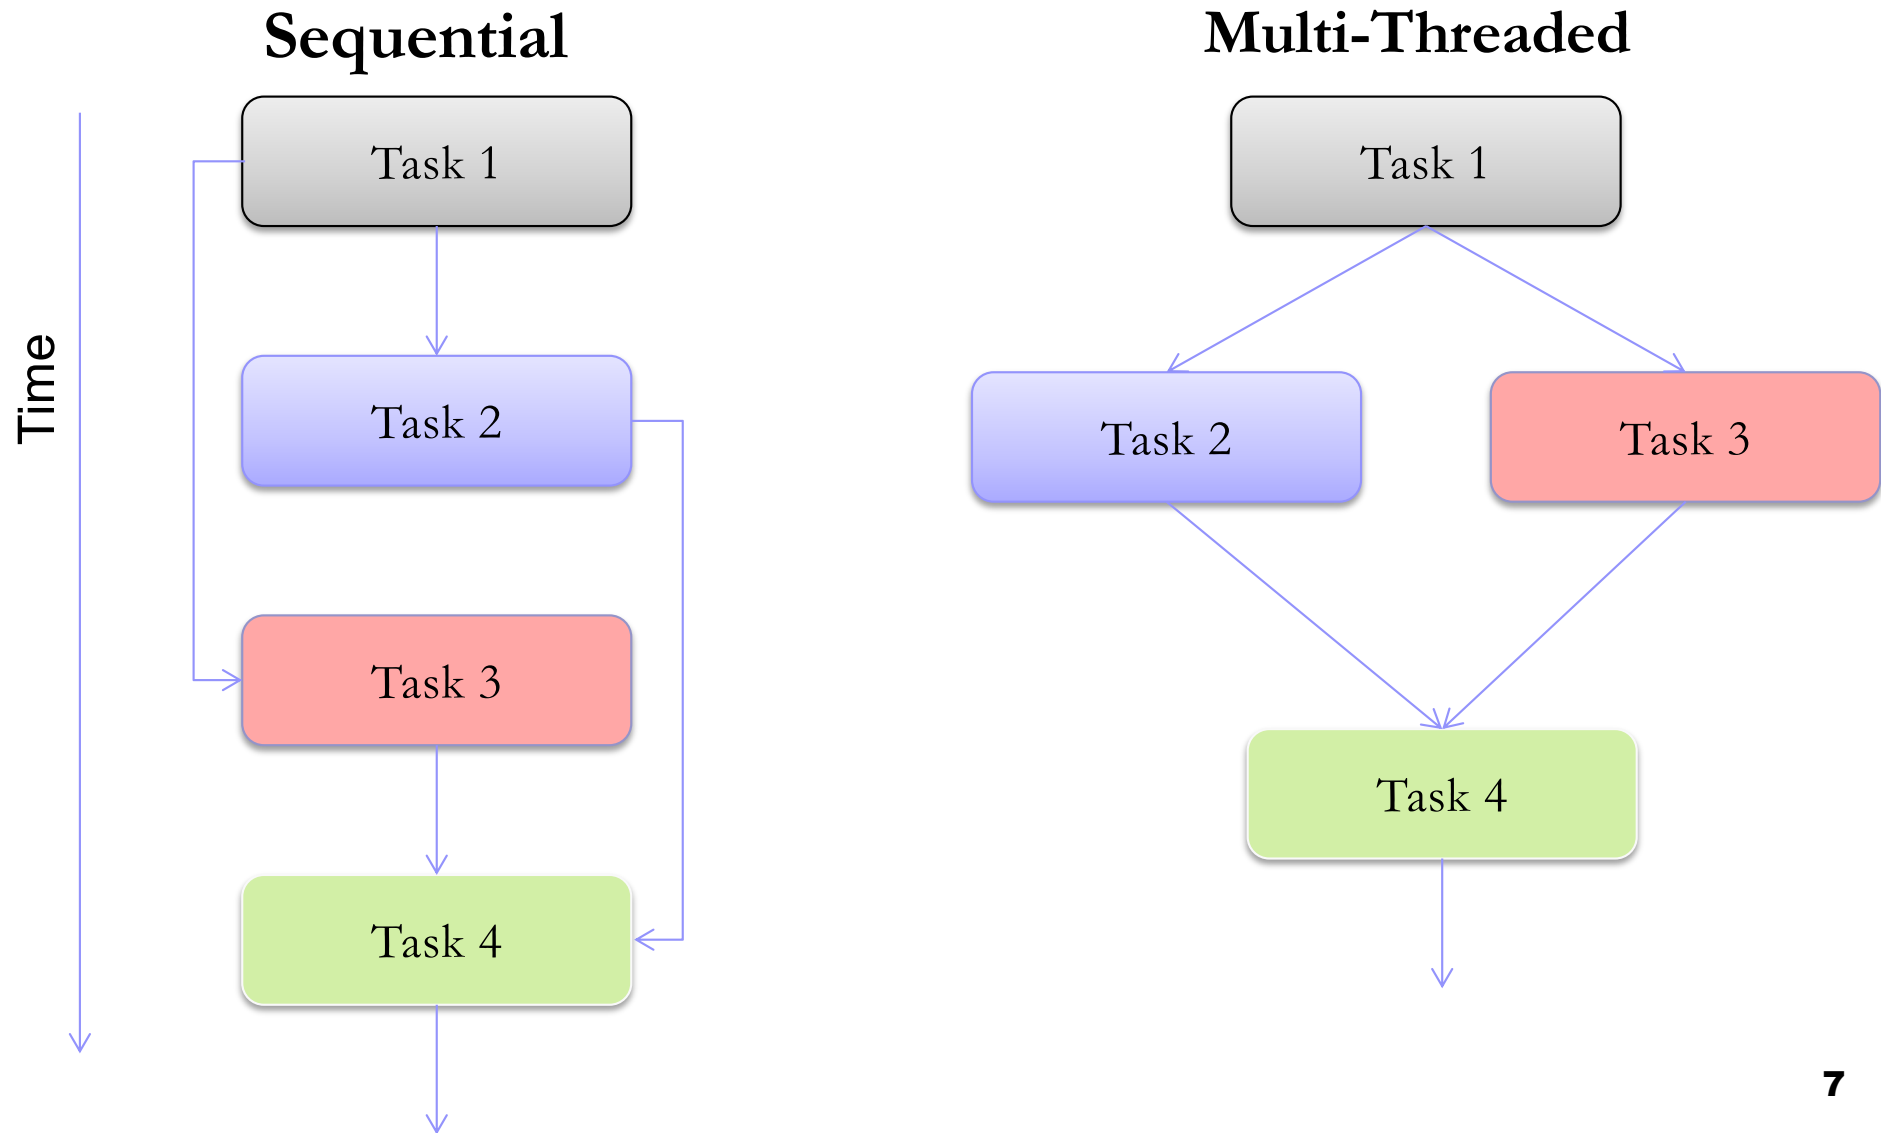


■ Memory Space



New things to worry about - Nondeterminism

- What might be the problem here?





New things to worry about - Liveness

- Previously we've worried about safety, i.e. writing our code so that it performs *correctly*
- With threads we now need to also worry about *liveness*
 - Not just concerned that code doesn't perform correctly - i.e. that bad things don't happen
 - Also concerned code performs **correctly in time** - i.e. that good things do happen eventually
 - This second requirement means we need to concern ourselves what level of *eventually* is acceptable...

Threads in Java

- Can be created in two ways in Java
 - Using (and extending) the Thread class
 - Using the Runnable interface
- Extending is simple to implement, but has drawbacks
- What might this be?

```
public class HelloThread extends Thread {  
  
    public void run () {  
        System.out.println ("Hello from a thread!");  
    }  
  
    public static void main ( String args []) {  
        (new HelloThread()).start();  
    }  
  
}
```

```
public class HelloRunnable implements Runnable {  
  
    public void run () {  
        System.out.println ("Hello from a thread!");  
    }  
  
    public static void main ( String args []) {  
        (new Thread (new HelloRunnable ())).start();  
    }  
  
}
```



Thread lifecycle methods

Package java.lang

```
public class Thread implements Runnable {  
    public void start();  
    public void run();  
    public void stop(); //deprecated (internal race condition)  
    public void resume(); // deprecated  
    public void suspend(); // deprecated  
    public static void sleep(long millis);  
    public static void sleep(long millis, int nanos);  
    public boolean isAlive();  
    public void interrupt();  
    public boolean isInterrupted();  
    public static boolean interrupted();  
    public void join();  
}
```

How do I stop a thread?

Setting a flag:

```
public void run () {  
    while (!done ) {  
        // ...  
    }  
}
```

- Loop will exit (and thread will stop) on the next iteration after done is set to true

Interrupting:

```
public void run() {  
    while (!isInterrupted()) {  
        // ...  
    }  
}
```

- As above, when `isInterrupted()` is called for this thread, it will stop
- Except that if the thread is sleeping / waiting will throw `InterruptedException` and immediately return from `run()`



Race conditions

- Data synchronisation between threads can be an issue
- When two threads try to access/change data at the same time it's known as a **race condition**
- Two volunteers for a game.

Race condition with threads

```
public class MyCounter {  
  
    private int count = 0;  
  
    public void addTwo() {  
        count = count + 2;  
    }  
  
    public void subtractTwo() {  
        count = count - 2;  
    }  
  
    public int countValue () {  
        return count ;  
    }  
  
}
```

- Thread 1: (addTwo) Retrieve counter.
- Thread 2: (subtractTwo) Retrieve counter.

- Thread 1: Change retrieved value; counter value is 2.
- Thread 2: Change retrieved value; result is -2.

- Thread 1: Store result in counter (counter is 2).
- Thread 2: Store result in counter (counter is now -2).

Synchronisation

- Simply add the synchronize keyword to the methods in the class
 - When one thread is executing a synchronized method for an object, all other threads that invoke synchronized methods for the **same object block suspend execution** until the first thread is done with the object.

```
public class MyCounter {  
    private int count = 0;  
    public synchronized void addTwo() {  
        count = count + 2;  
    }  
    public synchronized void subtractTwo() {  
        count = count - 2;  
    }  
    public synchronized int countValue() {  
        return count ;  
    }  
}
```

Statement synchronisation

- Also possible to synchronise a set of statements, not necessarily a whole method:

```
public void doSomeSynchronisedStuff() {  
    // everything out here before is not synchronized  
    synchronized (this){  
        // everything in here is synchronised  
        ...  
    }  
    // everything out here afterwards is not synchronised  
either  
}
```



Atomic actions

- In programming, atomic actions are self-contained, they cannot be stopped in the middle
 - They either happen or they don't
- Atomic access can help here:
 - Java primitive data types: byte, short, int, long, float, double, boolean, char
 - Reads and writes are atomic for most primitive variables (all types except long and double).
 - Reads and writes are atomic for *all* variables declared **volatile** (*including* long and double variables).



Summary

- Explored concurrency
 - Why concurrency is required
 - Multi-tasking vs multi-threading
- Threads as a mechanism for concurrency in Java
 - Liveness
 - Race conditions
 - Synchronization
 - Atomic access
- **Next time: more on threading**